

Cache Concepts and Performance

DSC 315: Computer Organization & Operating Systems

Dr. Laltu Sardar

School of Data Science,
Indian Institute of Science Education and Research Thiruvananthapuram (IISER TVM)



February 17, 2026



Memory hierarchy

1 Input-output and I/O System Design

Cache Write Policies

- Cache improves performance by hiding memory latency
- Write policy controls how updates propagate to main memory
- Tradeoff between performance, complexity, and correctness

Write-through

- Writes update cache and main memory simultaneously
- Memory always holds the latest value
- Simple and predictable behavior

Write-through: Pros and Cons

Advantages

- Strong consistency
- Simple hardware design
- Easier crash recovery

Disadvantages

- Higher write latency
- Increased memory traffic

Write-back Cache

Behavior

- Writes update only the cache
- Memory updated later on cache eviction
- Modified lines marked as dirty

Characteristics

- Faster writes
- Memory may be temporarily stale

Memory Stall

- A memory stall occurs when the CPU must wait for data or instructions
- Caused by cache misses, TLB misses, or memory ordering delays
- Pipeline stages become idle while waiting for memory
- Main memory latency can be hundreds of CPU cycles
- Memory stalls reduce IPC and overall performance

Intuition: the CPU pauses execution because the required data has not arrived yet.

Write-back: Pros and Cons

Advantages

- Low write latency
- Reduced memory bandwidth usage
- High overall performance

Disadvantages

- More complex hardware
- Risk of data loss on power failure
- Harder cache coherence in multicore systems

Summary and Use Cases


Aspect	Write-through	Write-back
Write latency	High	Low
Memory consistency	Always current	Can be stale
Hardware complexity	Low	High
Crash safety	Better	Worse
Performance	Lower	Higher

Typical use

- Write-through: embedded and safety-critical systems
- Write-back: modern CPUs and high-performance systems

Address (showing bit positions)

31 ... 14 13 ... 6 5 ... 2 1 0



One-way set associative**(direct mapped)**

Block	Tag	Data
0		
1		
2		
3		
4		
5		
6		
7		

Two-way set associative


Set	Tag	Data	Tag	Data
0				
1				
2				
3				



Four-way set associative


Set	Tag	Data	Tag	Data	Tag	Data	Tag	Data
0								
1								

Eight-way set associative (fully associative)

Tag	Data												







Why Cache Performance Matters

- Modern CPUs execute instructions much faster than main memory
- Cache misses cause the CPU to stall
- Overall performance depends heavily on how often and how long the CPU waits for memory

CPU Execution Time

CPU execution time can be expressed as

$$\text{CPU Time} = \text{CPU Clock Cycles} \times \text{Clock Cycle Time}$$

Total clock cycles include:

- CPU execution cycles
- Memory stall cycles

Memory-Stall Clock Cycles

Definition

Memory-stall clock cycles are cycles during which the CPU is stalled waiting for data or instructions from memory.

$$\text{Memory Stall Cycles} = \text{Read-Stall Cycles} + \text{Write-Stall Cycles}$$

Read-Stall Cycles

Read-stall cycles occur due to cache misses on read operations.

Sources of Read Stalls

- Instruction cache misses
- Data cache misses on loads

$$\text{Read-Stall Cycles} = \text{Read Misses} \times \text{Miss Penalty}$$

Write-Stall Cycles

Write-stall cycles depend on the cache write policy.

Possible Causes

- Write-through cache with a write buffer that becomes full
- Write misses that require fetching a block from memory

$$\text{Write-Stall Cycles} = \text{Write Misses} \times \text{Miss Penalty}$$

Instruction Miss Cycles

Instruction miss cycles occur when the instruction cache misses.

$$\text{Instruction Miss Cycles} = \text{Instruction Miss Rate} \times \text{Instructions} \times \text{Miss Penalty}$$

These stalls directly delay instruction fetch and pipeline progress.

Data Miss Cycles

Data miss cycles are caused by load and store instructions missing in the data cache.

$$\text{Data Miss Cycles} = \text{Data Miss Rate} \times \text{Memory Accesses} \times \text{Miss Penalty}$$

They often dominate performance in data-intensive programs.

Putting It All Together

Total Memory Stall Cycles

$$\text{Memory Stall Cycles} = \text{Instruction Miss Cycles} + \text{Data Miss Cycles}$$

Total CPU Clock Cycles

$$\text{CPU Clock Cycles} = \text{Execution Cycles} + \text{Memory Stall Cycles}$$

Average Memory Access Time

Definition

Average Memory Access Time (AMAT) measures the average time to access memory, including hits and misses.

$$\text{AMAT} = \text{Hit Time} + \text{Miss Rate} \times \text{Miss Penalty}$$

Improving Cache Performance

- Reduce miss rate
 - Larger cache
 - Better associativity
 - Improved block placement
- Reduce miss penalty
 - Multilevel caches
 - Critical-word-first and early restart
- Reduce hit time
 - Simple cache structures
 - Pipelined caches

Why Cache Misses Occur

- Cache misses increase memory stall cycles
- Reducing miss rate directly improves CPU performance
- Cache organization plays a major role in miss behavior

Fully Associative Cache

Key Idea

Any memory block can be placed in any cache line.

- No index field in the address
- Entire address tag is compared with all cache tags
- Lowest conflict miss rate
- Requires expensive hardware for parallel tag comparison

Set-Associative Cache

Key Idea

The cache is divided into sets, each containing multiple blocks.

- A block maps to exactly one set
- It can be placed in any line within that set
- Compromise between direct-mapped and fully associative caches

Locating a Block in the Cache

Address Breakdown

Address = Tag | Set Index | Block Offset

- Set index selects the set
- Tag comparison determines hit or miss
- Block offset selects the requested word or byte

Choosing Which Block to Replace

When a cache miss occurs and the set is full, one block must be replaced.

Replacement Policies

- Least Recently Used (LRU)
- First-In First-Out (FIFO)
- Random replacement

LRU generally gives the lowest miss rate but is more complex to implement.

Size of Tags vs Set Associativity

- Increasing associativity reduces conflict misses
- Higher associativity increases:
 - Tag storage per cache line
 - Hardware complexity for tag comparison
- Fully associative caches require the largest tag fields

Trade-off

Lower miss rate versus higher cost and longer hit time.

Reducing Miss Penalty with Multilevel Caches

Key Idea

Use multiple levels of cache with increasing size and access time.

- L1 cache: small, fast, low hit time
- L2 cache: larger, slower, lower miss rate
- L3 cache: even larger, shared among cores

Effect of Multilevel Caches

Effective Miss Penalty

$$\text{Effective Miss Penalty} = \text{L1 Miss Rate} \times \text{L2 Miss Penalty}$$

Each additional cache level significantly reduces the average memory access time.

Finally

- Associativity reduces conflict misses
- Replacement policies affect miss rate
- Higher associativity increases tag size and hardware cost
- Multilevel caches reduce effective miss penalty

Virtual Memory

Section 5.4

Motivation for Virtual Memory

- Programs may be larger than physical memory
- Multiprogramming requires memory isolation
- Efficient use of limited main memory is critical

Key Idea

Virtual memory gives each process the illusion of a large, private, contiguous memory space.

Virtual Address Space

- Each process has its own virtual address space
- Virtual addresses are mapped to physical addresses
- Only active portions of a program are kept in memory

Virtual Address \longrightarrow Physical Address

Paging

- Virtual memory is divided into fixed-size pages
- Physical memory is divided into page frames
- Page size is typically a power of 2

Terminology

- Page: unit of virtual memory
- Page frame: unit of physical memory

Virtual Address Format

For a page size of 2^n bytes:

$$\text{Virtual Address} = \text{Page Number} \mid \text{Page Offset}$$

- Page offset uses n bits
- Page number uses remaining bits

Page Table

- Page table maps virtual page numbers to physical frames
- Each process has its own page table

Page Table Entry

Typically contains:

- Frame number
- Valid or invalid bit
- Protection bits
- Dirty bit

Address Translation Using Page Table

- 1 Extract page number from virtual address
- 2 Use page number to index page table
- 3 Obtain physical frame number
- 4 Concatenate with page offset

Physical Address = Frame Number | Page Offset

Page Fault

Definition

A page fault occurs when a referenced page is not in physical memory.

- Operating system is invoked
- Required page is fetched from disk
- Page table is updated
- Instruction is restarted

Demand Paging

- Pages are loaded only when needed
- Reduces memory usage
- Increases efficiency for large programs

Key Property

Programs exhibit locality of reference.

Locality of Reference

- Temporal locality: recently used pages are reused
- Spatial locality: nearby addresses are accessed

Locality makes demand paging effective in practice.

Translation Lookaside Buffer (TLB)

- A small, fast cache for page table entries
- Stores recent virtual to physical translations

Goal

Reduce address translation time.

Effective Access Time with TLB

Let:

- t_{TLB} = TLB access time
- t_M = main memory access time
- h = TLB hit ratio

$$\text{EAT} = h \times (t_{TLB} + t_M) + (1 - h) \times (t_{TLB} + 2t_M)$$

Page Replacement

When memory is full, a page must be replaced.

Common Algorithms

- Least Recently Used (LRU)
- First-In First-Out (FIFO)
- Optimal replacement

Page Replacement Goal

- Minimize page fault rate
- Reduce disk I O operations
- Improve overall system performance

Effective Access Time with Page Faults

Let:

- p = page fault rate
- t_{PF} = page fault service time

$$\text{EAT} = (1 - p) \times t_M + p \times t_{PF}$$

Thrashing

Definition

Thrashing occurs when the system spends most of its time handling page faults.

- Insufficient physical memory
- Excessive multiprogramming
- Poor page replacement decisions

Summary

- Virtual memory separates address space from physical memory
- Paging enables efficient memory usage
- TLB improves address translation speed
- Page faults and replacement dominate performance

A Common Framework for Memory Hierarchies

Section 5.5

Parallelism and Private Caches

- Modern processors exploit parallelism using multiple cores
- Each core typically has private L1 instruction and data caches
- Private caches reduce latency and increase bandwidth
- However, multiple cached copies of the same memory block may exist

Key Challenge

Maintaining a consistent view of shared memory across all processors.

Shared Memory Multiprocessor Model

- All processors share a single physical address space
- Each processor executes its own instruction stream
- Memory can be accessed by any processor
- Caches replicate memory blocks for performance

Without coordination, replicated data can become inconsistent.

Definition of Cache Coherence

Cache Coherence

A system is cache coherent if it ensures that all processors see a consistent value for any shared memory location.

- A read returns the most recent write
- Writes are eventually visible to all processors
- All processors observe writes to the same location in the same order

The Cache Coherence Problem

- Processor P1 reads a memory block into its cache
- Processor P2 also reads the same block
- P1 modifies the block locally
- P2 still holds an outdated copy

Consequence

P2 may read stale data unless coherence is enforced.

Coherence Requirements

A cache coherence protocol must satisfy:

- **Write propagation:** a write by one processor eventually becomes visible to all
- **Write serialization:** writes to the same location are seen in the same order by all processors
- **Read correctness:** a read returns the value of the most recent write

Basic Approaches to Cache Coherence

- Snooping based protocols
- Directory based protocols

Key Difference

How coherence information is communicated among caches.

Snooping Based Coherence

- All caches are connected via a shared bus
- Each cache controller monitors all bus transactions
- Cache controllers take action when they observe relevant addresses
- Common in small scale multiprocessors

Broadcast based communication simplifies design but limits scalability.

Snooping Transactions

- Read miss generates a bus read request
- Write generates invalidation or update messages
- Other caches snoop the request and respond accordingly

Bus traffic grows rapidly as processor count increases.

Write Invalidate Protocol

- On a write, all other cached copies are invalidated
- Only one cache holds a writable copy
- Subsequent reads by other processors cause cache misses

Advantage

Lower bandwidth usage compared to write update.

Write Update Protocol

- On a write, updated data is broadcast to all caches
- All cached copies remain valid
- Readers always see up to date data

Disadvantage

High bus bandwidth consumption for frequent writes.

MESI Cache Coherence Protocol

MESI is a widely used write invalidate snooping protocol.

Cache Line States

- Modified
- Exclusive
- Shared
- Invalid

MESI State: Modified

- Cache line contains modified data
- Memory copy is stale
- This cache has exclusive ownership
- Data must be written back on eviction

MESI State: Exclusive and Shared

- Exclusive: block is clean and present in only one cache
- Shared: block is clean and may exist in multiple caches
- Reads are allowed in both states
- Writes require state transition

MESI State: Invalid

- Cache line does not contain valid data
- Any access results in a cache miss
- Line must be refetched from memory or another cache

Directory Based Coherence

- Used in large scale multiprocessors
- Each memory block has an associated directory entry
- Directory tracks which caches hold the block

Avoids broadcast traffic and improves scalability.

Directory Entry Structure

A directory entry typically contains:

- Block state
- Owner cache identifier
- Sharer list or bit vector

Directory based protocols trade storage for scalability.

False Sharing

- Occurs when processors modify different variables
- Variables reside in the same cache block
- Causes unnecessary invalidations

False sharing severely degrades parallel performance.

Reducing False Sharing

- Pad data structures to cache block boundaries
- Separate frequently written variables
- Improve memory layout for parallel programs

Performance Cost of Coherence

- Increased cache miss rate due to invalidations
- Additional memory latency from coherence traffic
- Reduced scalability with increasing core count

Efficient coherence is critical for parallel speedup.

Summary

- Cache coherence ensures correctness in shared memory systems
- Snooping and directory protocols enforce coherence
- MESI is a practical and widely deployed protocol
- False sharing is a major performance pitfall

Parallelism and Memory Hierarchies: Cache Coherence

Section 5.8



Parallelism and the Need for Cache Coherence

- Modern processors improve performance by executing multiple instruction streams in parallel
- Each processor core typically has private L1 caches to reduce memory access latency
- Caches replicate data blocks from main memory to improve locality and bandwidth
- Parallel programs frequently share data across cores

Fundamental Issue

Multiple cached copies of the same memory location can lead to inconsistent views of memory.

Shared Memory Multiprocessor Architecture

- All processors share a single physical address space
- Any processor can read or write any memory location
- Each processor has private caches holding copies of memory blocks
- Main memory acts as the backing store

Without a coordination mechanism, updates made by one processor may not be visible to others.

What Is Cache Coherence?

Cache Coherence

Cache coherence is the property that ensures all processors observe a consistent value for each shared memory location.

- A read must return the value of the most recent write
- Writes by one processor must eventually be seen by all others
- All processors must observe writes to the same address in the same order

Illustration of the Coherence Problem

- Processor P1 reads variable X into its cache
- Processor P2 also reads X into its cache
- P1 updates X locally in its cache
- P2 continues to read the old value of X

Result

Program correctness is violated unless coherence is enforced.

Coherence Requirements

Any cache coherence protocol must guarantee:

- **Write propagation:** updates eventually reach all caches
- **Write serialization:** all processors see writes in the same order
- **Read correctness:** a read returns the value of the most recent write

These properties are necessary for correct shared memory semantics.

Coherence vs Consistency

- Cache coherence concerns a single memory location
- Memory consistency concerns ordering of accesses to different locations
- A system can be coherent but still exhibit surprising behaviors without a strong consistency model

Cache coherence is a prerequisite for memory consistency.

Basic Approaches to Cache Coherence

- Snooping based coherence protocols
- Directory based coherence protocols

Key Distinction

How caches communicate coherence information to one another.

Snooping Based Cache Coherence

- All caches are connected to a shared communication medium
- Every cache controller monitors all memory transactions
- Requests are broadcast to all caches
- Caches take action based on observed requests

Snooping protocols are simple but do not scale well.

Snooping Protocol Operations

- Read miss generates a bus read request
- Write generates a bus write or invalidate request
- Other caches snoop the bus and update their state

Bus contention increases rapidly with processor count.

Write Invalidate Protocol

- On a write, all other cached copies are invalidated
- Only one cache holds a writable copy
- Subsequent reads by other processors cause cache misses

Benefit

Lower coherence traffic for multiple writes to the same block.

Write Update Protocol

- On a write, updated data is broadcast to all caches
- All cached copies remain valid and consistent
- Readers immediately see updated values

Drawback

High bandwidth consumption for write intensive workloads.

MESI Cache Coherence Protocol

- A write invalidate snooping protocol
- Widely implemented in commercial processors
- Each cache line has one of four states

States

Modified, Exclusive, Shared, Invalid

MESI State: Modified

- Cache line contains modified data
- Memory copy is stale
- Only one cache holds the block
- Write back required on eviction

MESI State: Exclusive

- Cache line is clean
- Present in exactly one cache
- Memory copy is up to date
- Write can proceed without invalidation

MESI State: Shared

- Cache line may be present in multiple caches
- All copies are clean
- Write requires invalidation of other copies

MESI State: Invalid

- Cache line does not contain valid data
- Any access results in a cache miss
- Block must be fetched from memory or another cache

Directory Based Cache Coherence

- Designed for large scale multiprocessors
- Avoids broadcast communication
- A directory keeps track of block ownership and sharers

Scales better than snooping based approaches.

Directory Entry Information

Each directory entry typically stores:

- State of the memory block
- Identity of the owner cache
- List or bitmap of sharer caches

Directory storage overhead grows with system size.

False Sharing

- Occurs when different processors update different variables
- Variables reside in the same cache block
- Causes repeated invalidations and cache misses

False sharing degrades performance without true data sharing.

Reducing False Sharing

- Align shared variables to cache block boundaries
- Pad frequently written variables
- Redesign data structures for parallel access

Performance Impact of Cache Coherence

- Increased memory access latency
- Higher cache miss rate due to invalidations
- Reduced scalability for write shared data

Efficient coherence protocols are critical for parallel speedup.

Summary

- Cache coherence maintains correctness in shared memory systems
- Snooping and directory protocols enforce coherence
- MESI is a practical and widely used protocol
- False sharing is a major performance bottleneck



THANK YOU

FOR YOUR ATTENTION

Dr. Laltu Sardar, Assistant Professor, IISER Thiruvananthapuram

laltu.sardar@iisertvm.ac.in, laltu.sardar.crypto@gmail.com

Course webpage: https://laltu-sardar.github.io/courses/corgos_2026.html.

-  Carl Hamacher, Zvonko Vranesic, and Safwat Zaky.
Computer Organization.
McGraw-Hill, 6th edition, 2012.
-  John L. Hennessy and David A. Patterson.
Computer Architecture: A Quantitative Approach.
Morgan Kaufmann, 6th edition, 2017.
-  Vincent P. Heuring and Harry F. Jordan.
Computer System Design and Architecture.
Pearson, 2nd edition, 2007.
-  David A. Patterson and John L. Hennessy.
Computer Organization and Design: The Hardware/Software Interface.
Morgan Kaufmann, 4th edition, 2014.
-  William Stallings.
Computer Organization and Architecture: Designing for Performance.
Pearson, 10th edition, 2016.